

Inter Process Communication



What to Expect?

- ★ W's of Inter Process Communication
- ★ Various IPC mechanisms in Linux
 - ◆ Pipes
 - ◆ FIFOs
 - ◆ Message Queues
 - ◆ Shared Memory & Process Semaphores



Why separate IPC mechanisms?

- * Processes are non-sharing entities
- * But, we may need communication
- * Already looked at one mechanism
 - ♦ Signals
- * But those are very basic
 - ♦ Mostly asynchronous
 - ♦ Connection-less
 - ♦ No data communication
- * Separate mechanisms are answer to all those



What is IPC?

- ★ Inter Process Communication
- ★ Mechanism whereby one process communicates with another process
- ★ Communication = Exchange of Data
- ★ Directions
 - ◆ Uni-directional
 - ◆ Bi-directional
 - ◆ Multi-directional



Example for IPC

- ★ Print the filenames in a directory
- ★ 'ls' outputs the filenames
- ★ 'lpr' prints its input
- ★ But ls & lpr are separate processes
- ★ What are the solutions?
 - ◆ Put into a file → ls > file; lpr < file
 - ◆ Setup an IPC, say a pipe (|) → ls | lpr



Types of IPC mechanisms

- ★ Pipe: Allows the flow of data in one direction only
- ★ Named Pipe: Pipe with a specific name
- ★ Message Queues: Message passing using a queue
- ★ Shared Memory & Semaphore
 - ◆ Shared Memory allows the interchange of data through a defined area of memory
 - ◆ Semaphore is used in solving problems associated with synchronization and avoiding race conditions during sharing
- ★ Mapped Memory: Similar to shared memory, but use file instead of memory
- ★ Why so many?
 - ◆ Based on the requirements, flexibility & benefits



Pipes



W's of a Pipe

- ★ Connects 2 ends to allow transfer
- ★ One-way Serial Communication Channel
- ★ One end: Output (write end)
- ★ Another end: Input (read end)
- ★ Limited Capacity
- ★ Provides automatic Synchronization
 - Read blocks on no data
 - Write blocks on full data



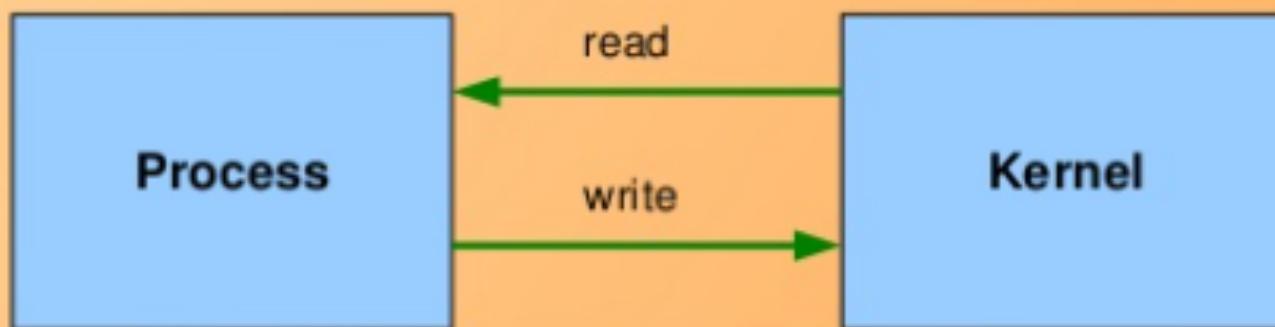
Pipe Usage

- * On Shell
 - ♦ Using |
- * In Programming / C
 - ♦ Creation
 - Using popen(), pclose()
 - Using system call pipe()
 - ♦ Communication as with any other fd
 - Using system calls read(), write(), ...



Pipe in a Process

- ★ Pipe Creation: Kernel sets up two file descriptors
 - 1st for Read: Path for obtaining the data
 - 2nd for Write: Path for data input
- ★ Pipe to communicate with itself !!!
 - May be among threads of a process



Making Sense with Pipes

- * Info: Child process inherits parent's open file descriptors
- * This Info + Pipe Creation = IPC within the family
- * Final Critical Decision
 - Remember pipes are half-duplex
 - So, who communicates with whom?



Multi-Process Communication

- * Usual system calls operating on file descriptors
- * Send data to the pipe: write()
- * Retrieve data from the pipe: read()
- * Though calls like lseek(), don't work



Pipes Extras & Limitations

- * Two-way Pipes
 - Open two pipes
 - Assign the file descriptors, appropriately
- * pipe() call must PRECEDE a call to fork()
- * Pipes are only for related processes
 - That is processes with a related ancestry
- * What about unrelated processes?
- * There are special type of pipes called ...



FIFOs



W's of a FIFO

- * A Pipe with a name in the file system
- * Also called Named Pipe
- * Properties identical to those of a Pipe
 - including the Serial Communication
 - which is now termed as First-In-First-Out
- * Except that now processes can be unrelated
- * Achieved by it being a device special file
 - which persists irrespective of Process existence
 - And any process can operate on it



FIFO Usage

- * On Shell
 - Creation: mknod, or mkfifo
 - Communication: cat, echo, ...
- * In Programming / C
 - Creation
 - System Call: mknod(), open(), close()
 - Library Function: mkfifo(), fopen(), fclose()
 - Communication
 - System Calls: read(), write(), ...
 - Library Functions: fread(), fwrite, fputs, ...



Experiment on Shell

- * mkfifo /tmp/fifo Open another shell
- * ls -l /tmp/fifo
- * cat < /tmp/fifo <Type some text>
- <See the text> * rm /tmp/fifo



Shared Memory

W's of Shared Memory

- * Rawest Communication Mechanism
- * Piece of Memory which is Shared
 - ♦ Between two or more processes
- * Fastest form of IPC
 - ♦ As no unnecessary copying of data
 - ♦ And no other overheads – not even synchronization :)
- * Synchronization is Users responsibility
 - ♦ As desired by him
 - ♦ Process Semaphores is the Mechanism provided



Linux Memory Model

- * All accesses in Linux are through virtual addresses
- * Each Process has its “own” Virtual Address Space
 - ♦ Which is split into Virtual Pages
- * Each Process maintains a mapping table
 - ♦ From its Physical Pages to these Virtual Pages
 - ♦ For access to its actual data
- * And mappings of multiple processes can point to the same physical page
 - ♦ Enabling the sharing of “actual” memory



Shared Memory Usage

* Shell Way

- Create: ipcmk -M <size> [-p <mode>]
- Get Information: ipcs [-m [-i id]]
- Remove: ipcrm [-M <shmkey> | -m <shmid>]

* Programming Way in C (using sys calls)

- Physical Memory Allocation: `shmget()`
- Mapping / Attaching to Virtual Memory: `shmat()`
- Access through usual (Virtual) Address Dereferencing
- Unmapping / Detaching from Virtual Memory: `shmdt()`
- Physical Memory Deallocation & Other Shared Memory Control Operations: `shmctl()`



Physical Memory Allocation

- * Done by only one of the sharing Processes
 - ♦ `int shmget(key_t key, size_t size, int shm_flg);`
- * Returns
 - ♦ Shared Memory id on success
 - ♦ -1 on error (and sets errno)
- * key: System-wide unique id for a Shared Memory
 - ♦ Same as in Message Queue
 - ♦ Process gets a particular shared memory id using its creation key
- * size: Number of bytes to allocate or get for the shared memory
- * `shm_flg`: operation | permissions
 - ♦ `IPC_CREATE, IPC_EXCL`



Shared Memory Attachment

Mapping to Process Virtual Address Space

★ Done by all the sharing Processes to Access the Shared Memory

- ◆ `void *shmat(int shm_id, const void *shm_addr, int shm_flg);`

★ Returns

- ◆ Virtual Start Address of the attached Shared Memory on success
- ◆ `MAP_FAILED` i.e. `(void *)(-1)` on error (and sets `errno`)

★ `shm_id`: Shared Memory id returned by `shmget()`

★ `shm_addr`

- ◆ Requested process virtual start address
- ◆ `NULL` enables Kernel to choose one

★ `shm_flg`

- ◆ `SHM_RND`: round down `shm_addr` to a multiple of the page size
- ◆ `SHM_RDONLY`: shared memory will be only read, not written

★ Children created by `fork()` inherit attached shared segments



Shared Memory Detachment

Unmapping from Process Virtual Address Space

- * Done by all the Processes, which did attach, or inherit
 - `int shmdt(const void *shm_addr);`
- * Returns: 0 on success; -1 on error (and sets `errno`)
- * `shm_addr`: Process virtual address map returned by `shmat()`
- * Call to `_exit()` or `exec*`() automatically detaches
- * Detachment may deallocate the shared memory
 - If done by the last process
 - And already marked to deallocate



Need for Synchronization

- * We already talked
 - ♦ Shared Memory has no Access Coordination
 - ♦ It has to be handled by us – the User
- * Moreover, now we observed
 - ♦ The issue with the Example
- * So, now its time to look into resolving it
- * That's where we go to the next topic ...



Specialized Semaphores

* Binary semaphore

- ♦ Semaphore controlling a single resource “availability”
- ♦ Possible values
 - 0 → Unavailable
 - 1 → Available
- ♦ Typically initialized to 0

* Mutex

- ♦ Specialized Binary Semaphore
- ♦ Used for a “Usable” rather than a “Consumable” Resource
- ♦ Process using it, only releases it
 - Giving a notion of ownership
- ♦ Typically initialized to 1

